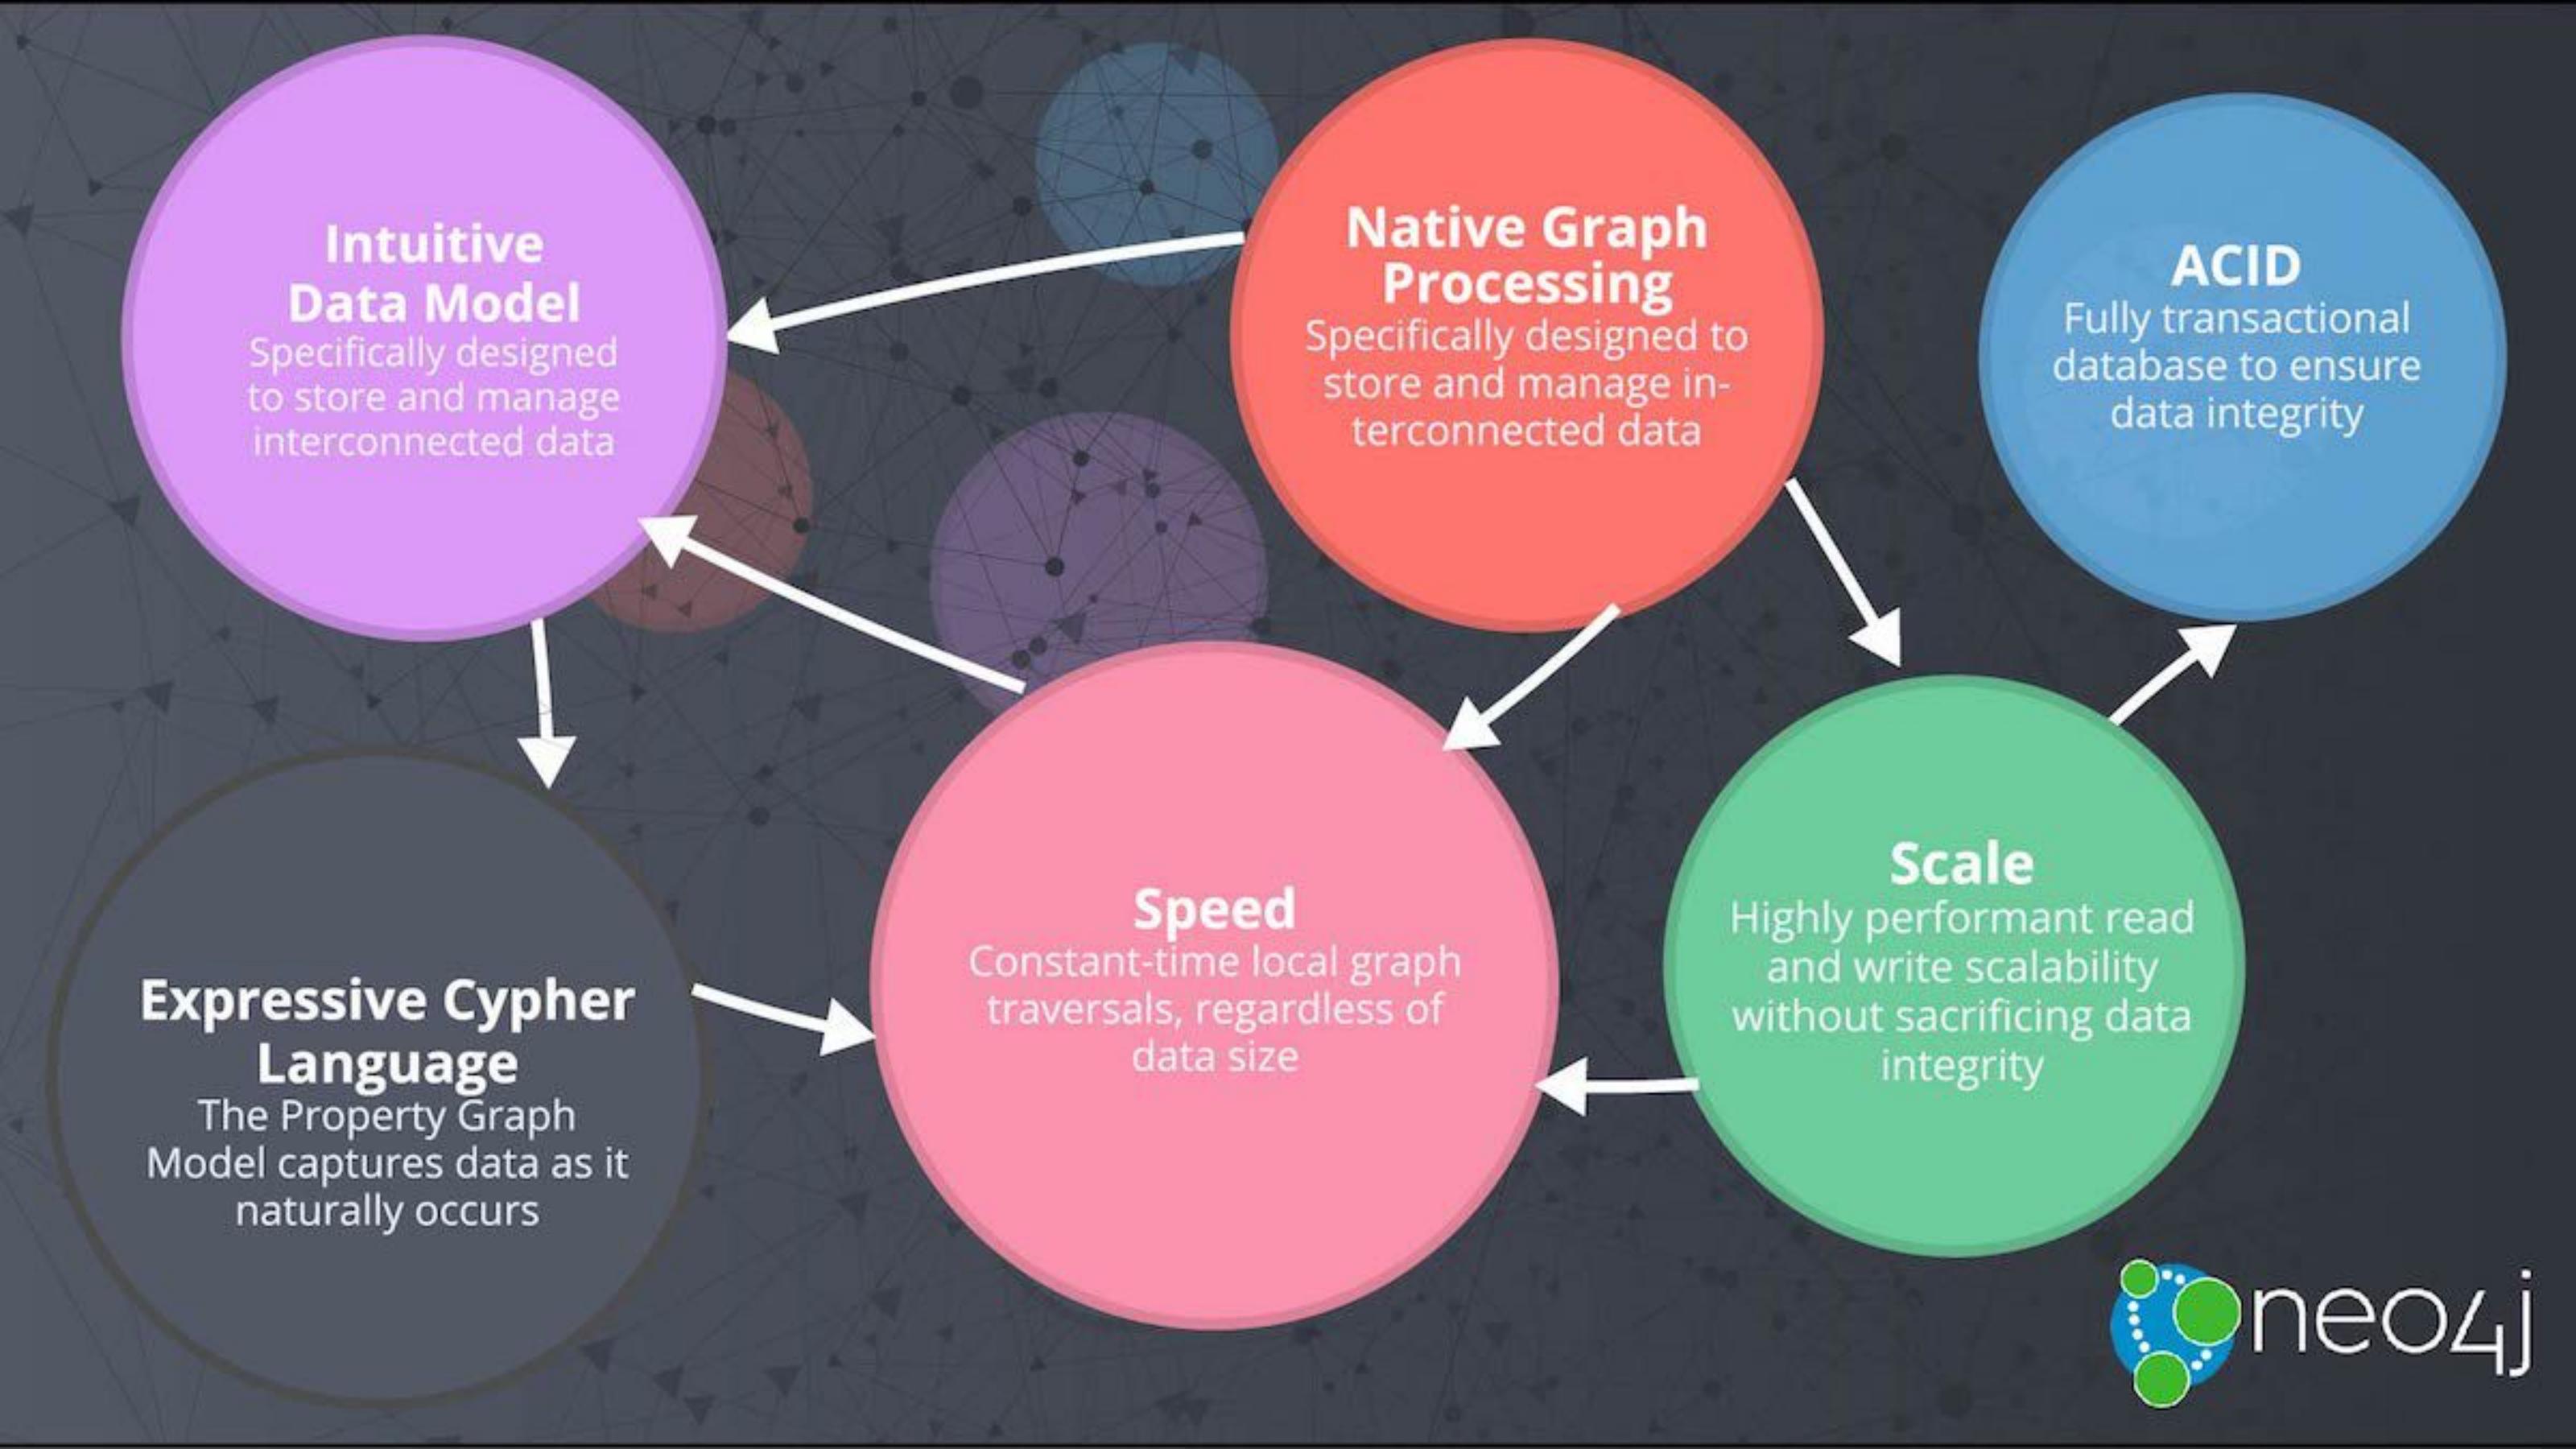


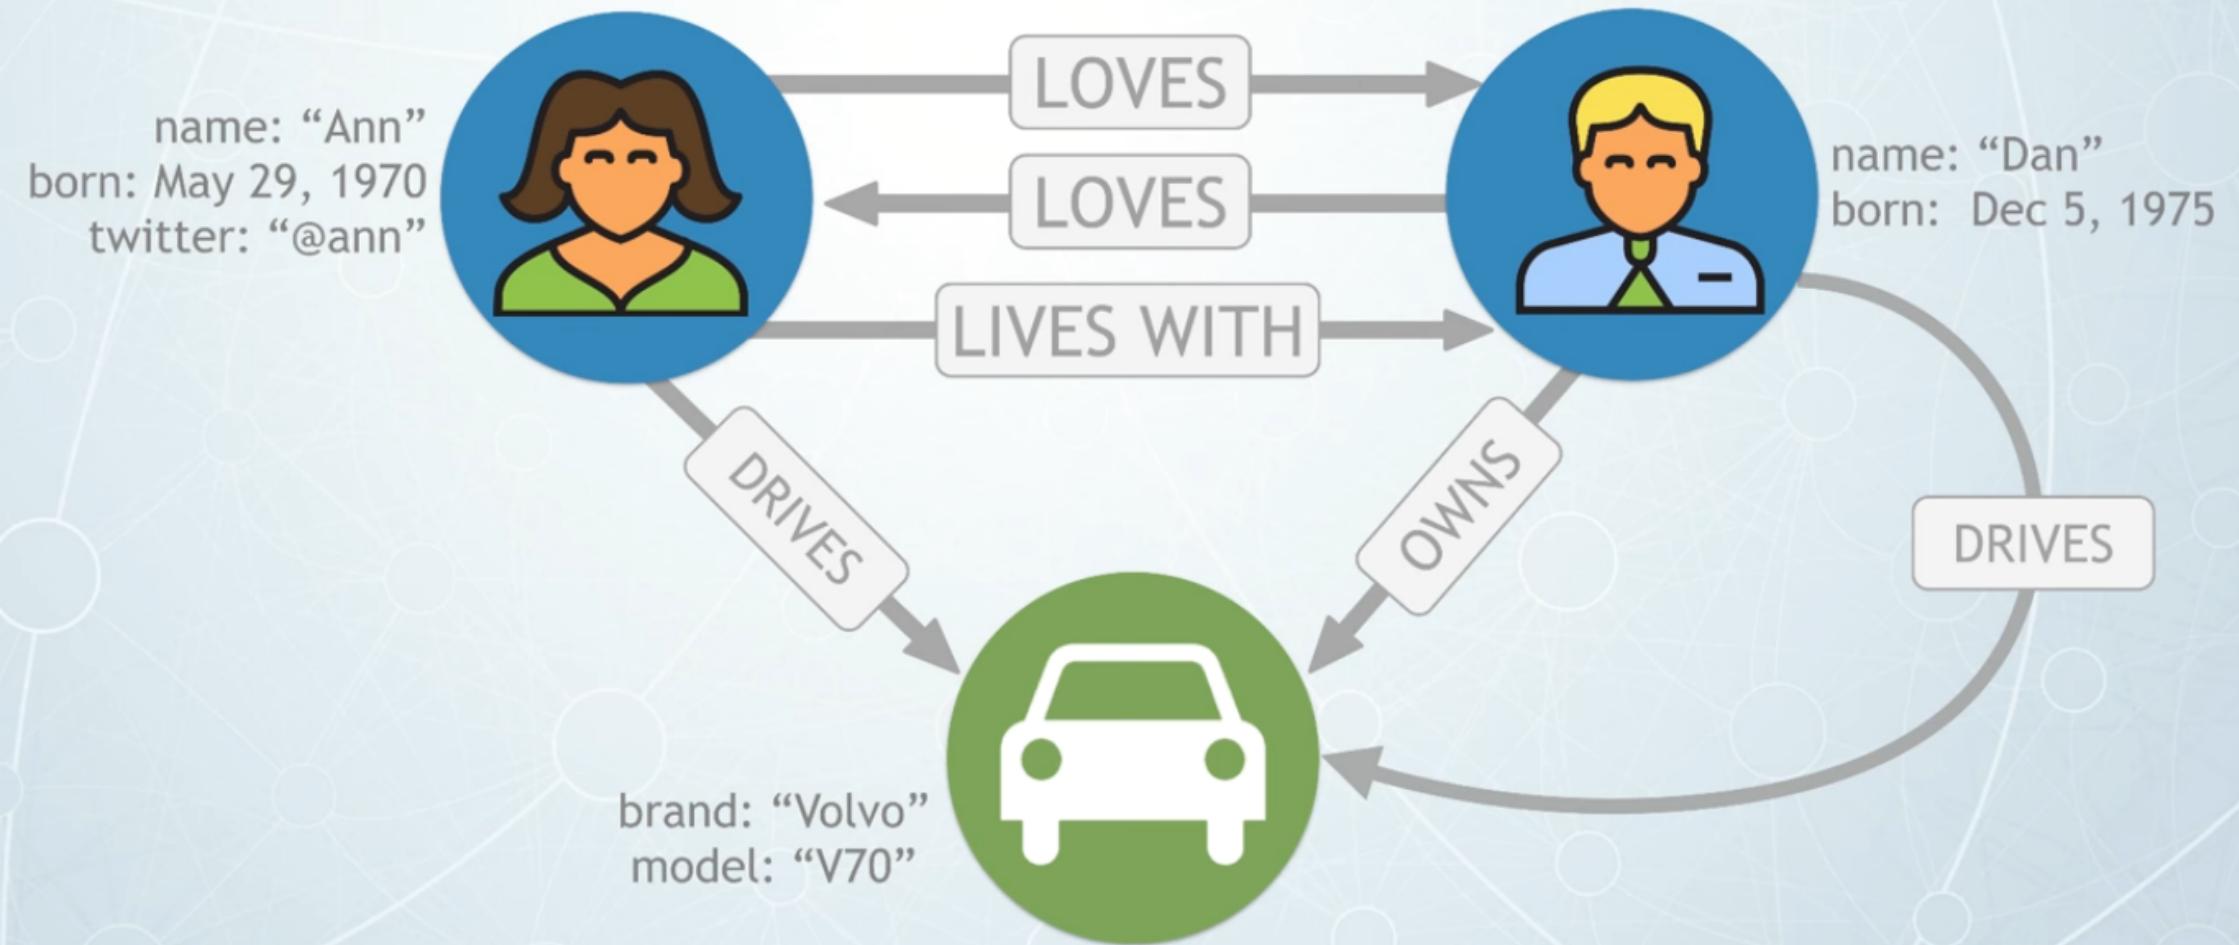


# Credits

- Most of the slides/images are taken from an online course, [Intro to Graph Databases Series](#), taught by Ryan Boyd (Neo4j Head of Developer Relations).
- Few are taken from the tutorials found on Neo4j's website.



# Detailed Property Graph



# **Basic Create and Query**

# Creating the Data



NODE



NODE

CREATE (:Person { name:“Ann”} ) - [:LOVES]-> (:Person { name:“Dan”} )

LABEL

PROPERTY

LABEL

PROPERTY

# **Patterns use AsciiArt**

# AsciiArt for Nodes

**Nodes are surrounded by parenthesis**

( ) or (p)

**Labels, or tags, start with : and group nodes by roles or types**

(p:Person:Mammal)

**Nodes can have properties**

(p:Person {name: 'Veronica'})

# AsciiArt for Relationships

**Relationships are wrapped with hyphens or square brackets**

--> or -[h:HIRED]->

**Direction of relationship is specified with < >**

(p1)-[:HIRED]->(p2)    Or    (p1)<-[:HIRED]-(p2)

**Relationships have properties too**

-[:HIRED {type: 'fulltime'}]->

# Patterns

**Patterns are drawn by connecting nodes and relationships with hyphens, optionally specifying a direction with > and < signs.**

○-[]-○

○-[]->○

○<-[]-○

# What were those?

**Relationships are wrapped with hyphens or square brackets**

--> or -[h:HIRED]->

**Direction of relationship is specified with < >**

(p1)-[:HIRED]->(p2)    Or    (p1)<-[:HIRED]-(p2)

**Relationships have properties too**

-[:HIRED {type: 'fulltime'}]->

# Variables

## Used to refer to Relationships

- [h:HIRED] ->

## Used to refer to Nodes

(p1)-[:HIRED]->(p2)    Or    (p1)<-[:HIRED]-(p2)

# How do I find Ann's Car?

```
MATCH  
  (:Person {name: 'Ann'})-[:DRIVES]->(c:Car)  
RETURN  
  c
```



# Is there another way?

```
MATCH  
  (a:Person)-[:DRIVES]->(c:Car)  
WHERE  
  a.name='Ann'  
RETURN  
  c
```



# Whom does Ann Love?



MATCH

(:Person {name:“Ann”})-[:LOVES]->(op:Person)

RETURN

op

# Whom does Ann Love?



MATCH

```
(:Person {name:"Ann"})-[:LOVES]->(op:Person)
```

RETURN

```
op
```

**It's Dan, of course!**



# Who drives a car owned by a lover?

```
MATCH
  (p1:Person)-[:DRIVES]->(c:Car)-[:OWNED_BY]->(p2:Person)<-[{:LOVES}]->(p1)
RETURN
  p1
```

# Components of a Cypher Query

```
MATCH path = (:Person)-[:ACTED_IN]->(:Movie)  
RETURN path
```

**MATCH** and **RETURN** are Cypher keywords

**path** is a variable

**:Movie** is a node label

**:ACTED\_IN** is a relationship

# **More on Syntax**

# Graph vs Tabular Results

```
1 MATCH (p:Person {name:"Tom Hanks"})-[r:ACTED_IN|DIRECTED]-(m:Movie)
2 RETURN p,r,m;
```

\$ MATCH (p:Person {name:"Tom Hanks"})-[r:ACTED\_IN|DIRECTED]-(m:Movie) RET...

Graph

Person(13) Movie(12)

ACTED\_IN(12) DIRECTED(1)

Displaying 13 nodes, 13 relationships.

```
1 MATCH (p:Person {name:"Tom Hanks"})-[r:ACTED_IN|DIRECTED]-(m:Movie)
2 RETURN p.name, type(r), m.title;
```

\$ MATCH (p:Person {name:"Tom Hanks"})-[r:ACTED\_IN|DIRECTED]-(m:Movie) RET...

Table

A Text

</> Code

p.name	type(r)	m.title
"Tom Hanks"	"ACTED_IN"	"Charlie Wilson's War"
"Tom Hanks"	"ACTED_IN"	"The Polar Express"
"Tom Hanks"	"ACTED_IN"	"A League of Their Own"
"Tom Hanks"	"ACTED_IN"	"Cast Away"
"Tom Hanks"	"ACTED_IN"	"Apollo 13"
"Tom Hanks"	"ACTED_IN"	"The Green Mile"
"Tom Hanks"	"ACTED_IN"	"The Da Vinci Code"
"Tom Hanks"	"ACTED_IN"	"Cloud Atlas"
"Tom Hanks"	"DIRECTED"	"That Thing You Do"
"Tom Hanks"	"ACTED_IN"	"That Thing You Do"
"Tom Hanks"	"ACTED_IN"	"Joe Versus the Volcano"
"Tom Hanks"	"ACTED_IN"	"Sleepless in Seattle"
"Tom Hanks"	"ACTED_IN"	"You've Got Mail"

Started streaming 13 records after 2 ms and completed after 5 ms.

# Graph vs Tabular Results

```
MATCH (p:Person {name:"Tom Hanks"})-[r:ACTED_IN|DIRECTED]-(m:Movie)  
RETURN p,r,m;
```



```
MATCH (p:Person {name:"Tom Hanks"})-  
[r:ACTED_IN|DIRECTED]-(m:Movie)  
RETURN p,r,m;
```

Displaying 13 nodes, 13 relationships.

```
MATCH (p:Person {name:"Tom Hanks"})-[r:ACTED_IN|DIRECTED]-(m:Movie)  
RETURN p.name, type(r), m.title;
```

p.name	type(r)	m.title
"Tom Hanks"	"ACTED_IN"	"Charlie Wilson's War"
"Tom Hanks"	"ACTED_IN"	"The Polar Express"
"Tom Hanks"	"ACTED_IN"	"A League of Their Own"
"Tom Hanks"	"ACTED_IN"	"Cast Away"
"Tom Hanks"	"ACTED_IN"	"Apollo 13"
"Tom Hanks"	"ACTED_IN"	"The Green Mile"
"Tom Hanks"	"ACTED_IN"	"The Da Vinci Code"
"Tom Hanks"	"ACTED_IN"	"Cloud Atlas"
"Tom Hanks"	"ACTED_IN"	"The Invention of Hugo Cabret"
"Tom Hanks"	"ACTED_IN"	"Inception"

```
MATCH (p:Person {name:"Tom Hanks"})-  
[r:ACTED_IN|DIRECTED]-(m:Movie)  
RETURN p.name, type(r), m.title;
```

Started streaming 13 records after 2 ms and completed after 5 ms.

# Tabular Results

Properties are accessed using: **{variable}.{property\_key}**

# Case Sensitivity

## **Case Sensitive**

Node labels

Relationship types

Property keys

## **Case Insensitive**

Cypher keywords

# Case Sensitivity

## **Case Sensitive**

:Person

:ACTED\_IN

name

## **Case Insensitive**

MaTcH

return

# **Tip for SETing properties**

# Using JSON Syntax for SET

```
MATCH
  (a:Person)-[:DRIVES]->(c:Car)
WHERE
  a.name='Ann'
SET
  c.brand='Volvo',
  c.model='V70'
RETURN
  c
```

# Using JSON Syntax for SET

Sometimes it's easiest to pass in an object to set properties

```
MATCH
  (a:Person)-[:DRIVES]->(c:Car)
WHERE
  a.name='Ann'
SET
  c.brand='Volvo',
  c.model='V70'
RETURN
  c
```

```
MATCH
  (a:Person)-[:DRIVES]->(c:Car)
WHERE
  a.name='Ann'
SET c += {brand:'Volvo', model:'V70'}
RETURN
  c
```

# **Ensuring Uniqueness**



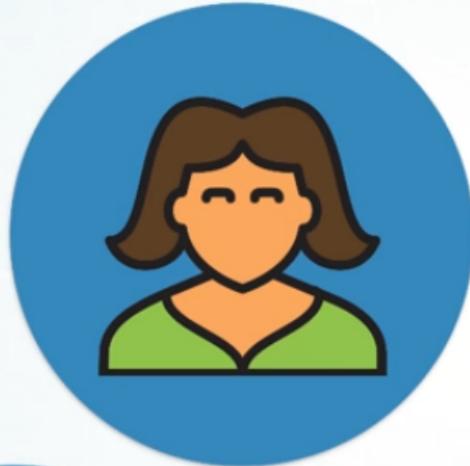
name: "Ann"



name: "Ann"



name: "Ann"



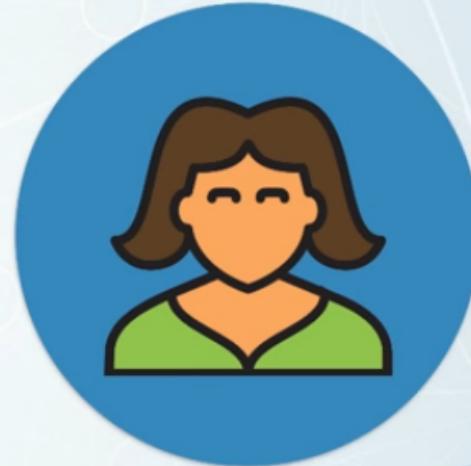
name: "Ann"



name: "Ann"



name: "Ann"



name: "Ann"



name: "Ann"



name: "Ann"



name: "Ann"



name: "Ann"

name: "Ann"

# There can be only One!

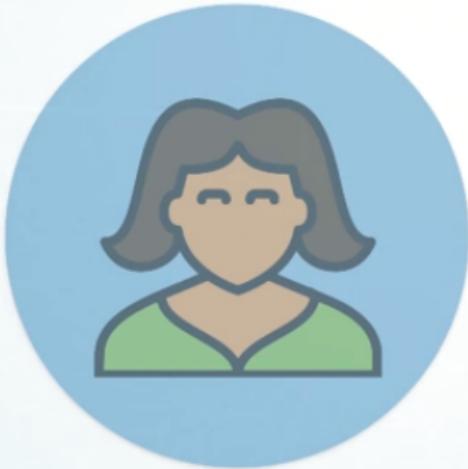
name: "Ann"



```
CREATE CONSTRAINT ON (p:Person)  
ASSERT p.name IS UNIQUE
```

# I want another Ann!!

name: "Ann"



name: "Ann"



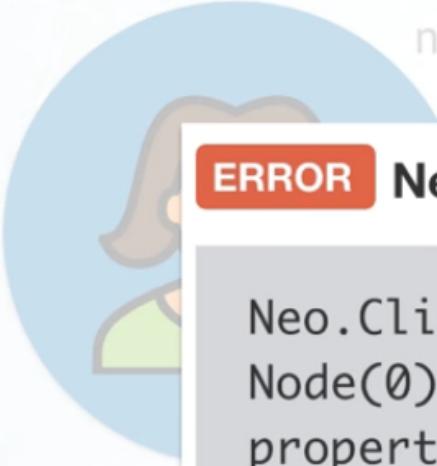
`CREATE (:Person {name:"Ann"})`

**ERROR** Neo.ClientError.Schema.ConstraintValidationFailed

Neo.ClientError.Schema.ConstraintValidationFailed:  
Node(0) already exists with label `Person` and  
property `name` = 'Ann'

# **Uniqueness on Create**

# I want Ann to Have a Pet Dog

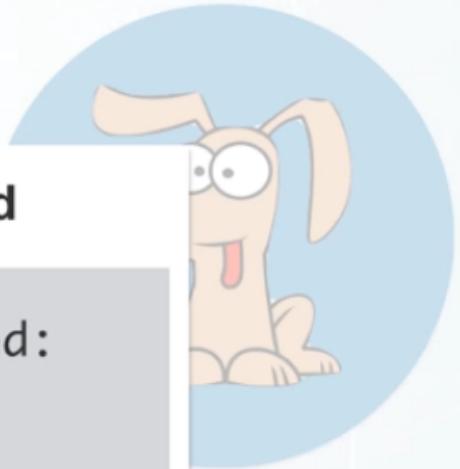


name: "Ann"

**ERROR**

**Neo.ClientError.Schema.ConstraintValidationFailed**

Neo.ClientError.Schema.ConstraintValidationFailed:  
Node(0) already exists with label `Person` and  
property `name` = 'Ann'



```
CREATE (a:Person {name:"Ann"})
```

```
CREATE (a)-[:HAS_PET]->(:Dog {name:"Sam"})
```

# I want Ann to Have a Pet Dog



```
MERGE (a:Person {name:"Ann"})
```

```
CREATE (a)-[:HAS_PET]->(:Dog {name:"Sam"})
```

# I want Ann to Have a Pet Dog



```
MERGE (a:Person {name:"Ann"})
```

```
ON CREATE SET
```

```
    a.twitter = "@ann"
```

```
MERGE (a)-[:HAS_PET]->(:Dog {name:"Sam"})
```

# **Let's Aggregate**

# Aggregates

Aggregate queries in Cypher are a little bit different than in SQL as we don't need to specify a grouping key.

~~GROUP BY p.name~~

# Aggregates

We implicitly group by any non-aggregate fields in the **RETURN** statement

```
// implicitly groups by p.name
MATCH (p:Person)-[:ACTED_IN]->(m:Movie)
RETURN p.name, count(*) AS numberOfMovies
```

# Aggregates

```
// implicitly groups by p.name
MATCH (p:Person)-[:ACTED_IN]->(m:Movie)
RETURN p.name, count(*) AS numberOfMovies
```

p.name	numberOfMovies
"River Phoenix"	1
"Halle Berry"	1
"Ben Miles"	3
"Ethan Hawke"	1

# Other Aggregating Functions

## Aggregating Functions

`count(*)`

The number of matching rows.

`count(variable)`

The number of non-null values.

`count(DISTINCT variable)`

All aggregating functions also take the DISTINCT operator, which removes duplicates from the values.

`collect(n.property)`

List from the values, ignores null.

`sum(n.property)`

Sum numerical values. Similar functions are `avg()`, `min()`, `max()`.

`percentileDisc(n.property, $percentile)`

Discrete percentile. Continuous percentile is

`percentileCont()`. The percentile argument is from 0.0 to 1.0.

`stDev(n.property)`

Standard deviation for a sample of a population. For an entire population use `stDevP()`.

[neo4j.com/docs/cypher-refcard/](https://neo4j.com/docs/cypher-refcard/)

# Neo4j is Extensible: Write your Own!

## 6.2.3.2. Writing a user-defined aggregation function

User-defined aggregation functions are annotated with `@UserAggregationFunction`. The annotated function must return an instance of an aggregator class. An aggregator class contains one method annotated with `@UserAggregationUpdate` and one method annotated with `@UserAggregationResult`. The method annotated with `@UserAggregationUpdate` will be called multiple times and allows the class to aggregate data. When the aggregation is done the method annotated with `@UserAggregationResult` is called once and the result of the aggregation will be returned. Valid output types are `long`, `Long`, `double`, `Double`, `boolean`, `Boolean`, `String`, `Node`, `Relationship`, `Path`, `Map<String, Object>`, or `List<T>`, where `T` can be any of the supported types.

For more details, see the API documentation for user-defined aggregation functions.

 The correct way to signal an error from within an aggregation function is to throw a `RuntimeException`.

```
package example;

import org.neo4j.procedure.Description;
import org.neo4j.procedure.Name;
import org.neo4j.procedure.UserAggregationFunction;
import org.neo4j.procedure.UserAggregationResult;
import org.neo4j.procedure.UserAggregationUpdate;

public class LongestString
{
    @UserAggregationFunction
    @Description("org.neo4j.function.example.longestString(string) - aggregates the longest string found")
    public LongStringAggregator longestString()
    {
        return new LongStringAggregator();
    }

    public static class LongStringAggregator
    {
        private int longest;
        private String longestString;

        @UserAggregationUpdate
        @Name("string")
        public void findLongest( String string )
        {
            if ( string != null && string.length() > longest )
            {
                longest = string.length();
                longestString = string;
            }
        }

        @UserAggregationResult
        public String result()
        {
            return longestString;
        }
    }
}
```

## User-defined aggregation functions

[neo4j.com/docs/](https://neo4j.com/docs/)

# Neo4j is Extensible: Use Open Source

## Aggregation Functions

apoc.agg.nth(value,offset)	returns non-null value of nth row (or -1 for last) offset is 0 based
apoc.agg.first(value)	returns first non-null value
apoc.agg.last(value)	returns last non-null value
apoc.agg.slice(value, start, length)	returns subset of non-null values, start is 0 based and length can be -1
apoc.agg.product(number)	returns given product for non-null values
apoc.agg.median(number)	returns median for non-null numeric values
apoc.agg.percentiles(value, [percentiles = 0.5,0.75,0.9,0.95,0.99])	returns given percentiles for integer values
apoc.agg.statistics(value, [percentiles = 0.5,0.75,0.9,0.95,0.99])	returns numeric statistics (percentiles, min,minNonZero,max,total,mean,stdev) for values

Aggregation Functions in APOC library

[neo4j-contrib.github.io/neo4j-apoc-procedures/](https://neo4j-contrib.github.io/neo4j-apoc-procedures/)

# **The WHERE clause**

## More on WHERE

```
MATCH (p:Person)-[r:ACTED_IN]->(m:Movie)
WHERE m.released >= 2000
RETURN m.released, a.name;
```

# Comparison Operators

=, <>, <, >, <=, >=, IS NULL, IS NOT NULL

# Note on null

**null** is not **null**

null represents missing or undefined values.

You do not store a null value in a property. It just doesn't exist on that particular node.

# Another Comparison Operator

$=\sim$

# Regular Expressions

```
MATCH (p:Person)-[r:ACTED_IN]->(m:Movie)
WHERE p.name =~ "K.+" OR m.released > 2000
RETURN p,r,m
```

# Boolean Logic

```
MATCH (p:Person)-[r:ACTED_IN]->(m:Movie)
WHERE p.name =~ "K.+" OR m.released > 2000
RETURN p,r,m
```

# Boolean Operators

**AND, OR, XOR, NOT**

# Grouping Boolean Operators

```
MATCH (p:Person)-[r:ACTED_IN]->(m:Movie)
WHERE m.released > 2000 OR
      (m.released = 1997 AND m.title='As Good as It Gets')
RETURN p.name, m.title, m.released
```